

# Efficient renaming in Conflict-free Replicated Data Types (CRDTs)

Case Study of a Sequence CRDT : LogootSplit

---

Matthieu Nicolas ([matthieu.nicolas@inria.fr](mailto:matthieu.nicolas@inria.fr))

COAST team

**Supervised by** Gérald Oster and Olivier Perrin

April 21, 2020

- Working on data replication in Local-first softwares<sup>[1]</sup>
- Currently focusing on *Sequence Conflict-free Replicated Data Types (CRDTs)*
  - Notably used for real-time collaborative text editing

---


<sup>[1]</sup>Martin Kleppmann et al. Local-first software: you own your data, in spite of the cloud. In *Proceedings of the 2019 ACM SIGPLAN International Symposium on New Ideas, New Paradigms, and Reflections on Programming and Software, Onward!* 2019, pages 154–178, Athens, Greece. Association for Computing Machinery, 2019. ISBN: 9781450369954. DOI: 10.1145/3359591.3359737. URL: <https://doi.org/10.1145/3359591.3359737> .

- State of the art of *Sequence CRDTs*
- Elements are ordered by their identifier, noted here with the following formalism:  $position_{offset}^{node\ id\ node\ seq}$

---

<sup>[2]</sup>Luc André et al. Supporting adaptable granularity of changes for massive-scale collaborative editing. In *International Conference on Collaborative Computing: Networking, Applications and Worksharing - CollaborateCom 2013*, pages 50–59, Austin, TX, USA. IEEE Computer Society, October 2013. DOI: 10.4108/icst.collaboratecom.2013.254123 .

- State of the art of *Sequence CRDTs*
- Elements are ordered by their identifier, noted here with the following formalism:  $position_{offset}^{node\ id\ node\ seq}$




**Figure 1:** State of a sequence which contains the elements "hlo" and their corresponding identifiers

---

[2] Luc André et al. Supporting adaptable granularity of changes for massive-scale collaborative editing. In *International Conference on Collaborative Computing: Networking, Applications and Worksharing - CollaborateCom 2013*, pages 50–59, Austin, TX, USA. IEEE Computer Society, October 2013. DOI: 10.4108/icst.collaboratecom.2013.254123 .

- State of the art of *Sequence CRDTs*
- Elements are ordered by their identifier, noted here with the following formalism:  $position_{offset}^{node\ id\ node\ seq}$



**Figure 1:** State of a sequence which contains the elements "hlo" and their corresponding identifiers




**Figure 2:** State of a sequence which contains the block "hlo"

---


[2] Luc André et al. Supporting adaptable granularity of changes for massive-scale collaborative editing. In *International Conference on Collaborative Computing: Networking, Applications and Worksharing - CollaborateCom 2013*, pages 50–59, Austin, TX, USA. IEEE Computer Society, October 2013. DOI: 10.4108/icst.collaboratecom.2013.254123 .

# Example




**Figure 3:** Example of concurrent *insert* operations

# Example



**Figure 3:** Example of concurrent *insert* operations


# Example



**Figure 3:** Example of concurrent *insert* operations



# Example




**Figure 3:** Example of concurrent *insert* operations

## Limits

- Unbounded size of identifiers
- Overhead of the data structure increasing over time

## Limits

- Unbounded size of identifiers
- Overhead of the data structure increasing over time



**Figure 4:** Footprint of the data structure

**How to reduce the overhead introduced by  
the data structure ?**

- Core-nebula approach<sup>[3]</sup>
  - Reassigns shorter identifiers to elements. . .
  - . . . but requires consensus
- LSEQ<sup>[4]</sup>
  - Set of strategies to reduce the growth of identifiers . . .
  - . . . but overhead still proportional to number of elements

---

<sup>[3]</sup>Marek Zawirski et al. Asynchronous rebalancing of a replicated tree. In *Conférence Française en Systèmes d'Exploitation (CFSE)*, page 12, Saint-Malo, France, May 2011. URL: <https://hal.inria.fr/hal-01248197> .


<sup>[4]</sup>Brice Nédelec et al. A scalable sequence encoding for collaborative editing. *Concurrency and Computation: Practice and Experience*:e4108. URL: <https://onlinelibrary.wiley.com/doi/abs/10.1002/cpe.4108>.

## Our approach

Reassign shorter identifiers and aggregate them into blocks in a fully distributed manner

# RenamableLogootSplit


- Propose *RenamableLogootSplit*, *LogootSplit* with a *rename* operation
- Can be perform without coordination
- Today, focus on scenario without concurrent *rename* operations



**Figure 5:** Example of renaming

# RenamableLogootSplit

- Propose *RenamableLogootSplit*, *LogootSplit* with a *rename* operation
- Can be performed without coordination
- Today, focus on scenario without concurrent *rename* operations



**Figure 5:** Example of renaming

- Generates a new identifier for the first element, based on its previous identifier



# RenamableLogootSplit

- Propose *RenamableLogootSplit*, *LogootSplit* with a *rename* operation
- Can be performed without coordination
- Today, focus on scenario without concurrent *rename* operations




**Figure 5:** Example of renaming

- Generates a new identifier for the first element, based on its previous identifier
- Then generates contiguous identifiers for all following elements

# RenamableLogootSplit

- Propose *RenamableLogootSplit*, *LogootSplit* with a *rename* operation
- Can be performed without coordination
- Today, focus on scenario without concurrent *rename* operations

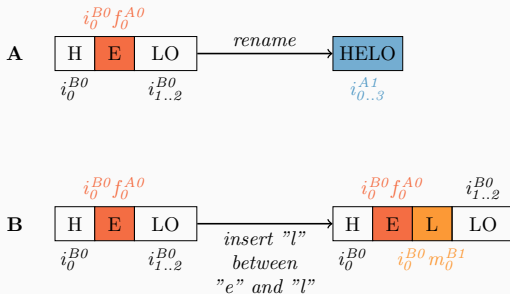


**Figure 5:** Example of renaming

- Generates a new identifier for the first element, based on its previous identifier
- Then generates contiguous identifiers for all following elements

# Handling concurrent operations - part 1

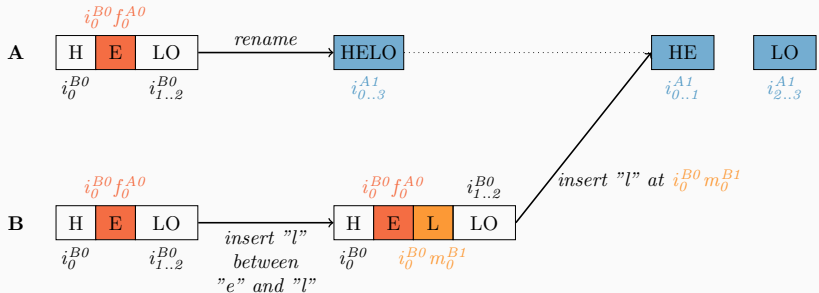
- Others may perform updates concurrently to a *rename* operation



**Figure 6:** Example of concurrent insert

# Handling concurrent operations - part 1


- Others may perform updates concurrently to a *rename* operation



**Figure 6:** Example of concurrent insert

# Handling concurrent operations - part 1

- Others may perform updates concurrently to a *rename* operation



**Figure 6:** Example of concurrent insert

- May lead to inconsistencies

# Handling concurrent operations - part 2

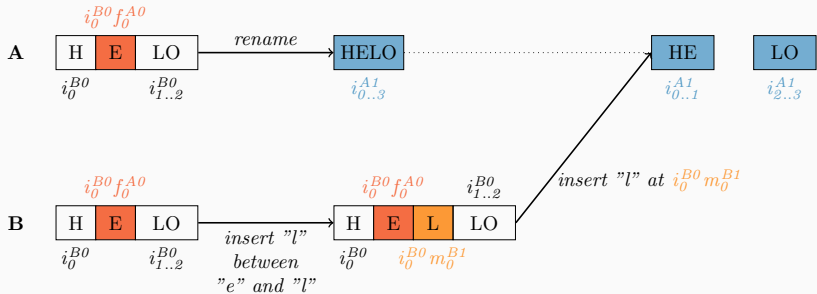



Figure 7: Example of concurrent insert


# Handling concurrent operations - part 2



**Figure 7:** Example of concurrent insert

- Use *epoch-based* system to track concurrent operations

# Handling concurrent operations - part 2




**Figure 7:** Example of concurrent insert

- Use *epoch-based* system to track concurrent operations
- Transform operations against *rename* ones (*OT*)




# Handling concurrent operations - part 2



**Figure 7:** Example of concurrent insert

- Use *epoch-based* system to track concurrent operations
- Transform operations against *rename* ones (*OT*)
  1. Find predecessor in former state ( $i_0^{B0} f_0^{A0}$ )


# Handling concurrent operations - part 2



**Figure 7:** Example of concurrent insert

- Use *epoch-based* system to track concurrent operations
- Transform operations against *rename* ones (OT)
  1. Find predecessor in former state ( $i_0^{B0} f_0^{A0}$ )
  2. Find its counterpart in new state ( $i_1^{A1}$ )

# Handling concurrent operations - part 2



**Figure 7:** Example of concurrent insert

- Use *epoch-based* system to track concurrent operations
- Transform operations against *rename* ones (*OT*)
  1. Find predecessor in former state ( $i_0^{B0} f_0^{A0}$ )
  2. Find its counterpart in new state ( $i_1^{A1}$ )
  3. Prepend it to given id to form new id ( $i_1^{A1} i_0^{B0} f_0^{A0}$ )

**Need to store former state until no more concurrent operations**

---

[5]Carlos Baquero et al. Making operation-based crdts operation-based. In Kostas Magoutis et al., editors, *Distributed Applications and Interoperable Systems*, pages 126–140, Berlin, Heidelberg. Springer Berlin Heidelberg, 2014 .

## Need to store former state until no more concurrent operations

- Can garbage collect it once the *rename* operation is causally stable<sup>[5]</sup>
- Can offload it to the disk meanwhile

---

<sup>[5]</sup>Carlos Baquero et al. Making operation-based crdts operation-based. In Kostas Magoutis et al., editors, *Distributed Applications and Interoperable Systems*, pages 126–140, Berlin, Heidelberg. Springer Berlin Heidelberg, 2014 .

## Need to store former state until no more concurrent operations

- Can garbage collect it once the *rename* operation is causally stable<sup>[5]</sup>
- Can offload it to the disk meanwhile

## Need to propagate former state to other nodes

---

<sup>[5]</sup>Carlos Baquero et al. Making operation-based crdts operation-based. In Kostas Magoutis et al., editors, *Distributed Applications and Interoperable Systems*, pages 126–140, Berlin, Heidelberg. Springer Berlin Heidelberg, 2014 .

## Need to store former state until no more concurrent operations

- Can garbage collect it once the *rename* operation is causally stable<sup>[5]</sup>
- Can offload it to the disk meanwhile

## Need to propagate former state to other nodes

- Can compress the operation to minimise bandwidth consumption

---

<sup>[5]</sup>Carlos Baquero et al. Making operation-based crdts operation-based. In Kostas Magoutis et al., editors, *Distributed Applications and Interoperable Systems*, pages 126–140, Berlin, Heidelberg. Springer Berlin Heidelberg, 2014 .

# Evaluation



# Scenario

- **Assumption:** Only one node can issue *rename* operations

## Ran simulations to evaluate proposed approach:

- 10 nodes share and edit a document collaboratively
- Nodes use either LogootSplit or RenamableLogootSplit according to session

# Scenario

- **Assumption:** Only one node can issue *rename* operations

## Ran simulations to evaluate proposed approach:

- 10 nodes share and edit a document collaboratively
- Nodes use either LogootSplit or RenamableLogootSplit according to session
- **Phase 1 (content generation):** 80/20% of *insert/remove*
- **Phase 2 (refactoring):** 50/50% of *insert/remove*
- Nodes switch to phase 2 when document reaches critical size (15 pages - 60k elements)

# Scenario

- **Assumption:** Only one node can issue *rename* operations

## Ran simulations to evaluate proposed approach:

- 10 nodes share and edit a document collaboratively
- Nodes use either LogootSplit or RenamableLogootSplit according to session
- **Phase 1 (content generation):** 80/20% of *insert/remove*
- **Phase 2 (refactoring):** 50/50% of *insert/remove*
- Nodes switch to phase 2 when document reaches critical size (15 pages - 60k elements)
- Overall, nodes perform 150k operations on the document
- In the case of *RenamableLogootSplit*, trigger a *rename* operation every 30k operations


## Results - Convergence

- Compared final content of nodes per sessions
- Did not observe any divergence

# Results - Convergence


- Compared final content of nodes per sessions
- Did not observe any divergence
- Empirical result, not a proof...
- ... but represents first step towards the validation

# Results - Overhead of the data structure



**Figure 8:** Evolution of the size of the document


# Results - Overhead of the data structure



**Figure 8:** Evolution of the size of the document

- *Rename* resets the overhead of the CRDT, if can garbage collect

# Results - Overhead of the data structure




**Figure 8:** Evolution of the size of the document


- *Rename* resets the overhead of the CRDT, if can garbage collect
- *Rename* still reduces by 66% the size otherwise



# Results - Integration time of insert operations




(a) Local operations




(b) Remote operations

Figure 9: Evolution of the integration time of *insert* operations

# Results - Integration time of insert operations



(a) Local operations




(b) Remote operations


**Figure 9:** Evolution of the integration time of *insert* operations

- *Rename* resets integration times of future operations

# Results - Integration time of insert operations



(a) Local operations




(b) Remote operations

**Figure 9:** Evolution of the integration time of *insert* operations


- *Rename* resets integration times of future operations
- Transforming concurrent operations is actually faster than applying them on former state

## Results - Integration time of rename operations



**Figure 10:** Evolution of the integration time of *rename* operations


## Results - Integration time of rename operations



**Figure 10:** Evolution of the integration time of *rename* operations

- Noticeable by users if delayed too much

## Results - Integration time of rename operations



**Figure 10:** Evolution of the integration time of *rename* operations

- Noticeable by users if delayed too much
- When to trigger *rename* operations?

# To wrap up

## Done

- Designed a *rename* operation for LogootSplit
- Defined transformation functions to deal with concurrent updates

---

[6]Matthieu Nicolas et al. MUTE: A Peer-to-Peer Web-based Real-time Collaborative Editor. In Proceedings of European Conference on Computer-Supported Cooperative Work - Panels, Posters and Demos, 2017 .

# To wrap up

## Done

- Designed a *rename* operation for LogootSplit
- Defined transformation functions to deal with concurrent updates

## Work in progress

- Implementing support of concurrent *rename* operations in MUTE<sup>[6]</sup>, our P2P collaborative text editor
- Proceeding to its validation
- Designing the strategy to trigger *rename* operations

---

[6]Matthieu Nicolas et al. MUTE: A Peer-to-Peer Web-based Real-time Collaborative Editor. In Proceedings of European Conference on Computer-Supported Cooperative Work - Panels, Posters and Demos, 2017 .



# To wrap up

## Done

- Designed a *rename* operation for LogootSplit
- Defined transformation functions to deal with concurrent updates

## Work in progress

- Implementing support of concurrent *rename* operations in MUTE<sup>[6]</sup>, our P2P collaborative text editor
- Proceeding to its validation
- Designing the strategy to trigger *rename* operations

## To do

- Prove formally the correctness of the mechanism
- Investigate the combination of OT techniques and CRDTs

---

<sup>[6]</sup>Matthieu Nicolas et al. MUTE: A Peer-to-Peer Web-based Real-time Collaborative Editor. In Proceedings of European Conference on Computer-Supported Cooperative Work - Panels, Posters and Demos, 2017 .

Thanks for your attention, any questions?



# LogootSplit identifiers

- To comply with these constraints, LogootSplit proposes identifiers composed of quadruplets of integers of the following form:

$\langle \textit{priority}, \textit{siteld}, \textit{seq}, \textit{offset} \rangle$

- *priority* allows to determine the position of this identifier compared to others
- *siteld* refers to the node's identifier, assumed to be unique
- *seq* refers to the node's logical clock, which increases monotonically with local operations
- *offset* refers to the element position in its original block

# Identifier constraints

- To fulfill their role, identifiers have to comply to several constraints:

## **Globally unique**

- Identifiers should never be generated twice, neither by different users nor by the same one at different times

## **Totally ordered**

- We should always be able to compare and order two elements using their identifiers

## **Dense set**

- We should always be able to add a new element, and thus a new identifier, between two others

*The topic of a later contribution*

# Handling concurrent rename

*The topic of a later contribution*

**rename operation not commutative**

# Handling concurrent rename

*The topic of a later contribution*

**rename operation not commutative**

**To fix this:**

- Define a total order between *rename* operations
- Pick a "winner" operation between concurrent *renames*
- Define additional transformation functions to *undo* the effect of "losing" ones

## Propose a strategy to avoid conflicting rename operations

- How to minimise likelihood of concurrent *rename* operations without coordinating?



## **Propose a strategy to avoid conflicting rename operations**

- How to minimise likelihood of concurrent *rename* operations without coordinating?

## **Propose a smarter strategy to choose the "winning" renaming**

- How to minimise the overall computations?

# Experimental settings

- Use Node.js version 13.1.0
- Obtained documents sizes using our fork of *object-sizeof* <sup>[7]</sup>
- Ran benchmarks on a workstation equipped of a Intel Xeon CPU E5-1620 (10MB Cache, 3.50 GHz) with 16GB of RAM running Fedora 31
- Measured times using `process.hrtime.bigint()`

---

<sup>[7]</sup><https://www.npmjs.com/package/object-sizeof>