Efficient renaming in Conflict-free Replicated Data Types (CRDTs)

Case Study of a Sequence CRDT : LogootSplit

**Matthieu Nicolas** (matthieu.nicolas@inria.fr), Gérald Oster, Olivier Perrin
COAST team

April 28, 2020

Context

- Working on data replication in Local-first softwares[1]
- Focusing on *Sequence Conflict-free Replicated Data Types (CRDTs)*
  - Real-time collaborative text editing

[1]Martin Kleppmann et al. Local-first software: you own your data, in spite of the cloud. In *Proceedings of the 2019 ACM SIGPLAN International Symposium on New Ideas, New Paradigms, and Reflections on Programming and Software*, Onward! 2019, pages 154–178, Athens, Greece. Association for Computing Machinery, 2019. ISBN: 9781450369954. DOI: 10.1145/3359591.3359737. URL: https://doi.org/10.1145/3359591.3359737 .

Conflict-free Replicated Data Types (CRDTs)[2]

- Replicated data structure

[2]Marc Shapiro et al. Conflict-free replicated data types. In *Proceedings of the 13th International Symposium on Stabilization, Safety, and Security of Distributed Systems*, SSS 2011, pages 386–400, 2011. DOI: 10.1007/978-3-642-24550-3_29.

Conflict-free Replicated Data Types (CRDTs)[2]

- Replicated data structure
- Updates performed without coordination

[2]Marc Shapiro et al. Conflict-free replicated data types. In *Proceedings of the 13th International Symposium on Stabilization, Safety, and Security of Distributed Systems*, SSS 2011, pages 386–400, 2011. DOI: 10.1007/978-3-642-24550-3_29 .

Conflict-free Replicated Data Types (CRDTs)[2]

- Replicated data structure
- Updates performed without coordination

[2]Marc Shapiro et al. Conflict-free replicated data types. In *Proceedings of the 13th International Symposium on Stabilization, Safety, and Security of Distributed Systems*, SSS 2011, pages 386–400, 2011. DOI: 10.1007/978-3-642-24550-3_29 .

Conflict-free Replicated Data Types (CRDTs)[2]

- Replicated data structure
- Updates performed without coordination

[2]Marc Shapiro et al. Conflict-free replicated data types. In *Proceedings of the 13th International Symposium on Stabilization, Safety, and Security of Distributed Systems*, SSS 2011, pages 386–400, 2011. DOI: 10.1007/978-3-642-24550-3_29 .

Conflict-free Replicated Data Types (CRDTs)[2]

- Replicated data structure
- Updates performed without coordination
- Strong Eventual Consistency

[2]Marc Shapiro et al. Conflict-free replicated data types. In *Proceedings of the 13th International Symposium on Stabilization, Safety, and Security of Distributed Systems*, SSS 2011, pages 386–400, 2011. DOI: 10.1007/978-3-642-24550-3_29 .

LogootSplit[3]

- State of the art of *Sequence CRDTs*
- Elements are ordered by their identifier, noted here with the following formalism: $position_{offset}^{node\_id\ node\_seq}$

_____

[3]Luc André et al. Supporting adaptable granularity of changes for massive-scale collaborative editing. In *International Conference on Collaborative Computing: Networking, Applications and Worksharing - CollaborateCom 2013*, pages 50–59, Austin, TX, USA. IEEE Computer Society, October 2013. DOI: 10.4108/icst.collaboratecom.2013.254123 .

LogootSplit[3]

- State of the art of *Sequence CRDTs*
- Elements are ordered by their identifier, noted here with the following formalism: $position_{offset}^{node\_id\ node\_seq}$

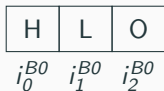| H	L	O

i_0^{B0} i_1^{B0} i_2^{B0}

Figure 1: State of a sequence which contains the elements "hlo" and their corresponding identifiers

[3]Luc André et al. Supporting adaptable granularity of changes for massive-scale collaborative editing. In *International Conference on Collaborative Computing: Networking, Applications and Worksharing - CollaborateCom 2013*, pages 50–59, Austin, TX, USA. IEEE Computer Society, October 2013. DOI: 10.4108/icst.collaboratecom.2013.254123 .

LogootSplit[3]

- State of the art of *Sequence CRDTs*
- Elements are ordered by their identifier, noted here with the following formalism: $position_{offset}^{node\_id\ node\_seq}$

H	L	O

i_0^{B0} i_1^{B0} i_2^{B0}

Figure 1: State of a sequence which contains the elements "hlo" and their corresponding identifiers

HLO

$i_{0..2}^{B0}$

Figure 2: State of a sequence which contains the block "hlo"

[3]Luc André et al. Supporting adaptable granularity of changes for massive-scale collaborative editing. In *International Conference on Collaborative Computing: Networking, Applications and Worksharing - CollaborateCom 2013*, pages 50–59, Austin, TX, USA. IEEE Computer Society, October 2013. DOI: 10.4108/icst.collaboratecom.2013.254123 .

4

LogootSplit[3]

- State of the art of *Sequence CRDTs*
- Elements are ordered by their identifier, noted here with the following formalism: $position_{offset}^{node\_id\ node\_seq}$


Figure 1: State of a sequence which contains the elements "hlo" and their corresponding identifiers

**Figure 2:** State of a sequence which contains the block "hlo"

Figure 3: State of a sequence which contains the elements "hlo" and their corresponding identifiers

[3]Luc André et al. Supporting adaptable granularity of changes for massive-scale collaborative editing. In *International Conference on Collaborative Computing: Networking, Applications and Worksharing - CollaborateCom 2013*, pages 50–59, Austin, TX, USA. IEEE Computer Society, October 2013. DOI: 10.4108/icst.collaboratecom.2013.254123 .

- **Evergrowing overhead:** impacts memory, bandwidth and CPU

Figure 4: Memory footprint of the data structure

- **Operation count:** 100k
- **Size of content:** 100KB
- **Size of data structure:** 20MB

How to reduce the overhead introduced by the data structure?

Our approach

Reassign shorter identifiers and aggregate them into blocks in a fully distributed manner

RenamableLogootSplit

- Propose *RenamableLogootSplit*, *LogootSplit* with a *rename* operation
- Can be perform without coordination
- In this talk, focus on scenario without concurrent *rename* operations

Figure 5: Example of renaming

RenamableLogootSplit

- Propose *RenamableLogootSplit*, *LogootSplit* with a *rename* operation
- Can be perform without coordination
- In this talk, focus on scenario without concurrent *rename* operations

Figure 5: Example of renaming

- Generates a new identifier for the first element, based on its previous identifier

RenamableLogootSplit

- Propose *RenamableLogootSplit*, *LogootSplit* with a *rename* operation
- Can be perform without coordination
- In this talk, focus on scenario without concurrent *rename* operations

Figure 5: Example of renaming

- Generates a new identifier for the first element, based on its previous identifier
- Then generates contiguous identifiers for all following elements

RenamableLogootSplit

- Propose *RenamableLogootSplit*, *LogootSplit* with a *rename* operation
- Can be perform without coordination
- In this talk, focus on scenario without concurrent *rename* operations

Figure 5: Example of renaming

- Generates a new identifier for the first element, based on its previous identifier
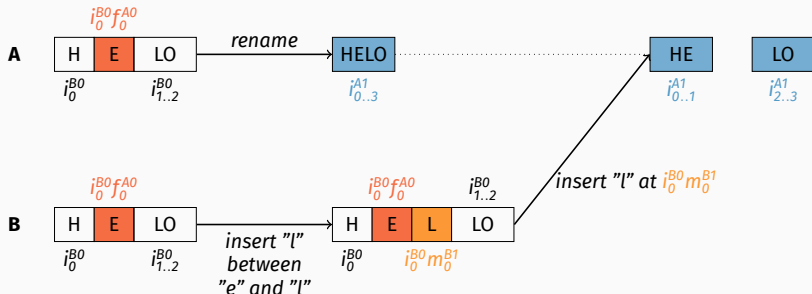- Then generates contiguous identifiers for all following elements

Figure 6: Example of concurrent insert

Figure 6: Example of concurrent insert

- Use *epoch-based* system to track concurrent operations

Figure 6: Example of concurrent insert

- Use *epoch-based* system to track concurrent operations
- Transform operations against *rename* ones (*OT*)

Handling concurrent operations

Figure 6: Example of concurrent insert

- Use *epoch-based* system to track concurrent operations
- Transform operations against *rename* ones (OT)
  1. Find predecessor in former state ($i_0^{B0} f_0^{A0}$)
  2. Find its counterpart in new state (i_1^{A1})
  3. Prepend it to given id to form new id ($i_1^{A1} i_0^{B0} f_0^{A0}$)

7

# Handling concurrent operations

Figure 6: Example of concurrent insert

- Use *epoch-based* system to track concurrent operations
- Transform operations against *rename* ones (OT)
  1. Find predecessor in former state ($i_0^{B0} f_0^{A0}$)
  2. Find its counterpart in new state (i_1^{A1})
  3. Prepend it to given id to form new id ($i_1^{A1} i_0^{B0} f_0^{A0}$)

Downsides

Need to store former state until no more concurrent operations

- Can garbage collect it once the *rename* operation is causally stable[4]
- Can offload it to the disk meanwhile

Need to propagate former state to other nodes

- Can compress the operation to minimise bandwidth consumption
- Can trigger *rename* operations at a given number of blocks

---

[4]Carlos Baquero et al. Making operation-based crdts operation-based. In Kostas Magoutis et al., editors, *Distributed Applications and Interoperable Systems*, pages 126–140, Berlin, Heidelberg. Springer Berlin Heidelberg, 2014 .

Evaluation

Ran simulations to compare RenamableLogootSplit performances to LogootSplit ones

Results - Convergence

- Compared final content of nodes per sessions
- Did not observe any divergence

- Empirical result, not a proof...
- ... but represents first step towards the validation

Results - Memory footprint

- **Phase 1 (content generation):** 80/20% of *insert*/*remove*
- **Phase 2 (editing):** 50/50% of *insert*/*remove*
- Nodes switch to phase 2 when document reaches critical size (15 pages - 60k elements)

Figure 7: Evolution of the size of the document

Results - Memory footprint

- **Phase 1 (content generation):** 80/20% of *insert*/*remove*
- **Phase 2 (editing):** 50/50% of *insert*/*remove*
- Nodes switch to phase 2 when document reaches critical size (15 pages - 60k elements)

Figure 7: Evolution of the size of the document

- *Rename* resets the overhead of the CRDT, if can garbage collect

Results - Memory footprint

- **Phase 1 (content generation):** 80/20% of *insert*/*remove*
- **Phase 2 (editing):** 50/50% of *insert*/*remove*
- Nodes switch to phase 2 when document reaches critical size (15 pages - 60k elements)
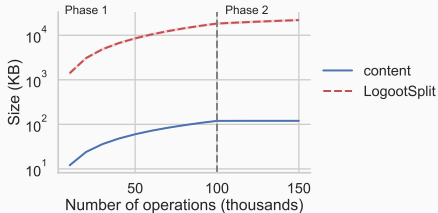

Figure 7: Evolution of the size of the document
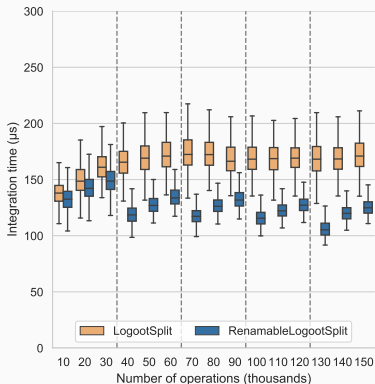
- *Rename* resets the overhead of the CRDT, if can garbage collect
- *Rename* still reduces by 66% the size otherwise

(a) Local operations

(b) Remote operations

Figure 8: Evolution of the integration time of *insert* operations

- *Rename* resets integration times of future operations
- Transforming concurrent operations is actually faster than applying them on former state

Figure 9: Evolution of the integration time of *rename* operations

- Noticeable by users if delayed too much

Research trail: propose strategies to retain acceptable integration time for *rename* operations

Conclusion

Done

- Designed a *rename* operation for LogootSplit
- Defined transformation functions to deal with concurrent updates

[5]Matthieu Nicolas et al. MUTE: A Peer-to-Peer Web-based Real-time Collaborative Editor. In Proceedings of European Conference on Computer-Supported Cooperative Work - Panels, Posters and Demos, 2017 .

Conclusion

Done

- Designed a *rename* operation for LogootSplit
- Defined transformation functions to deal with concurrent updates

Work in progress

- Implementing support of concurrent *rename* operations in MUTE[5], our P2P collaborative text editor
- Proceeding to its validation
- Designing strategies to trigger *rename* operations

[5]Matthieu Nicolas et al. MUTE: A Peer-to-Peer Web-based Real-time Collaborative Editor. In Proceedings of European Conference on Computer-Supported Cooperative Work - Panels, Posters and Demos, 2017 .

Conclusion

Done

- Designed a *rename* operation for LogootSplit
- Defined transformation functions to deal with concurrent updates

Work in progress

- Implementing support of concurrent *rename* operations in MUTE[5], our P2P collaborative text editor
- Proceeding to its validation
- Designing strategies to trigger *rename* operations

To do

- Prove formally the correctness of the mechanism
- Investigate the combination of OT techniques and CRDTs

[5]Matthieu Nicolas et al. MUTE: A Peer-to-Peer Web-based Real-time Collaborative Editor. In Proceedings of European Conference on Computer-Supported Cooperative Work - Panels, Posters and Demos, 2017 .

Thanks for your attention, any questions?

LogootSplit identifiers

- To comply with these constraints, LogootSplit proposes identifiers composed of quadruplets of integers of the following form:

$$position^{node\_id\ node\_seq}_{offset}$$

- *position* allows to determine the position of this identifier compared to others
- *node_id* refers to the node's identifier, assumed to be unique
- *node_seq* refers to the node's logical clock, which increases monotonically with local operations
- *offset* refers to the element position in its original block

Identifier constraints

- To fulfill their role, identifiers have to comply to several constraints:

Globally unique

- Identifiers should never be generated twice, neither by different users nor by the same one at different times

Totally ordered

- We should always be able to compare and order two elements using their identifiers

Dense set

- We should always be able to add a new element, and thus a new identifier, between two others

Related work

- Core-nebula approach[6]
    - Reassigns shorter identifiers to elements. . .
    - . . . but requires consensus
- LSEQ[7]
    - Set of strategies to reduce the growth of identifiers . . .
    - . . . but overhead still proportional to number of elements

[6]Marek Zawirski et al. Asynchronous rebalancing of a replicated tree. In *Conférence Française en Systèmes d'Exploitation (CFSE)*, page 12, Saint-Malo, France, May 2011. URL: https://hal.inria.fr/hal-01248197 .

[7]Brice Nédelec et al. A scalable sequence encoding for collaborative editing. *Concurrency and Computation: Practice and Experience*:e4108. URL: https://onlinelibrary.wiley.com/doi/abs/10.1002/cpe.4108.

Handling concurrent rename

The topic of a later contribution

rename operation not commutative

To fix this:

- Define a total order between *rename* operations
- Pick a "winner" operation between concurrent *renames*
- Define additional transformation functions to *undo* the effect of "losing" ones

Perspectives

Propose a strategy to avoid conflicting rename operations

- How to minimise likelihood of concurrent *rename* operations without coordinating?

Propose a smarter strategy to choose the "winning" renaming

- How to minimise the overall computations?

Experimental settings

- Use Node.js version 13.1.0
- Obtained documents sizes using our fork of *object-sizeof* [8]
- Ran benchmarks on a workstation equipped of a Intel Xeon CPU E5-1620 (10MB Cache, 3.50 GHz) with 16GB of RAM running Fedora 31
- Measured times using `process.hrtime.bigint()`

---

[8]https://www.npmjs.com/package/object-sizeof